

CERN CMS High-Level Trigger System: A Scientific High Performance Data Analytics Use Case

One of the most important High Performance Data Analytics systems in science is the real-time event filtering system at CERN's Large Hadron Collider (LHC), specifically the Compact Muon Solenoid (CMS) experiment's trigger system. The CMS High-Level Trigger (HLT) is a critical HPDA application that processes collision data in real-time to identify potentially important physics events among billions of interactions occurring every second. When the LHC operates at peak luminosity, approximately one billion proton-proton interactions occur every second at CMS, generating unprecedented data volumes that must be filtered in real-time to select only the most scientifically valuable events. The trigger system must reduce the initial data rate from 40 terabytes per second down to approximately 1 gigabyte per second suitable for permanent storage, making millisecond-level processing decisions on physics events to determine what gets recorded for analysis and what is discarded forever. The result is that CMS detects fundamental particles including the Higgs boson and searches for physics beyond the Standard Model, discoveries that would have been impossible without this sophisticated real-time data analysis capability.

Core Infrastructure & Scale

The CMS experiment operates at extraordinary scale in terms of collision events and data throughput. The LHC produces 40 million bunch crossings per second, with each bunch crossing producing approximately 1 megabyte of detector data. The system experiences a total collision rate of 33 megahertz at LHC Run 3 conditions with luminosity between $2\text{--}2.4 \times 10^{34} \text{ cm}^2 \text{ s}^{-1}$, creating 1.4 billion proton-proton interactions per second on average. Across the four LHC experiments, the entire facility generates approximately 20 petabytes of data annually that requires analysis and storage. At full operational capacity, CMS accepts approximately 110 kilohertz of events through the Level-1 hardware trigger and processes them through the High-Level Trigger software system, ultimately writing approximately 7.5 kilohertz to permanent storage. The CMS detector maintains over 360 to 400 compute nodes in its HLT processing farm, containing 30,512 CPU cores total and representing 7.2×10^5 HEPSPEC06 units of computational power. The detector processes approximately 1,200 algorithm instances from roughly 200 distinct algorithms on every event that passes to the HLT, with each event undergoing approximately 450 milliseconds of average processing time to make the keep/discard decision. The Worldwide LHC Computing Grid (WLCG) supporting all LHC experiments spans 1.4 million computer cores and 1.5

exabytes of storage distributed across over 170 computing sites in 42 countries, enabling analysis of LHC data by over 10,000 physicists worldwide.

Real-Time Processing Performance

The CMS trigger achieves extraordinary real-time processing performance through a two-stage architecture. The Level-1 trigger, implemented in custom hardware electronics, makes initial filtering decisions with a latency of only 3.8 microseconds, reducing the 40 megahertz input collision rate to 110 kilohertz. This hardware stage can only access simplified detector information from muon systems and calorimeters without tracker data due to latency constraints. The High-Level Trigger then takes these 110 kilohertz events and applies full detector information and sophisticated reconstruction algorithms, reducing the rate to approximately 7.5 kilohertz for disk storage. Each event requires approximately 450 milliseconds of processing time on average at the HLT stage. The data rate after Level-1 filtering reaches 9 gigabytes per second, which the HLT reduces to approximately 1 gigabyte per second for archival. The system currently runs at approximately 40 terabytes per second input data throughput. Modern developments include GPU acceleration integration to improve HLT processing speed and capability.

Technical Architecture

The CMS trigger system employs a sophisticated distributed computing architecture designed for real-time physics event selection. The Level-1 trigger consists of custom-built electronics designed specifically for particle detection, capable of making simple yes/no decisions about event interest based on calorimeter and muon chamber information within microsecond timescales. These decisions reduce the collision rate by a factor of approximately 400. The High-Level Trigger layer consists of a commercial PC farm comprising more than 1,000 standard computers working in parallel, with the HLT event filter farm containing approximately 360-400 compute nodes totaling 30,512 CPU cores. Recent implementations have begun incorporating graphics processing units alongside traditional CPUs to accelerate the most computationally intensive reconstruction algorithms. The trigger system must process structured numerical detector data from multiple heterogeneous sources including silicon pixel and strip tracker hits, electromagnetic calorimeter crystal energy deposits, hadron calorimeter tower signals, and muon chamber coordinates. Full event reconstruction becomes available only after the Level-1 trigger decision, at which point the complete event data approximately 1.2 megabytes per event at pileup conditions of 64 interactions is available to the HLT algorithms. The HLT applies approximately 1,200 algorithm

instances from roughly 200 distinct physics selection algorithms to decide whether an event merits permanent storage.

Machine Learning Algorithms and Physics Algorithms

The CMS HLT employs both traditional physics-based algorithms and increasingly sophisticated machine learning approaches. The system uses algorithms including Kalman filtering for track reconstruction, identifying particles from silicon tracker measurements with typical efficiency of 85-95% despite detector imperfections such as disabled pixel modules at levels reaching 5%. The HLT implements complex particle-flow reconstruction, combining information from all detector subsystems to create a complete picture of particle interactions. Modern trigger implementations include machine learning-based b-quark jet tagging algorithms that identify jets containing b-quarks with high efficiency, improving the discovery potential for Higgs boson decays and top quark events. The trigger employs regression algorithms for energy estimation in calorimeters and classification networks for differentiating signal physics processes from background processes.

Workflow Characteristics

The CMS trigger operates as a continuous real-time system processing event streams at unprecedented rates. The workflow characteristics include immediate trigger decisions made within 3.8 microseconds at Level-1 and approximately 450 milliseconds average at the High-Level Trigger, with no possibility for re-analysis of rejected events. The system runs 24/7 during LHC operational periods, making binary accept/reject decisions on events with permanent consequences. Trigger algorithms are continuously refined and updated during data-taking periods to respond to changing detector conditions, beam properties, and evolving physics priorities. Individual physicists and physics groups submit trigger algorithms and selections to be incorporated into the official trigger menu, which can contain hundreds of individual trigger paths searching for different physics signatures.

How the 5V Big Data Challenges Apply

Volume: The CMS experiment generates 20 petabytes of data annually across all LHC experiments. At peak luminosity, CMS alone processes 40 terabytes per second of detector data through the Level-1 trigger. Individual events can be 1.2 megabytes at high pileup, creating enormous aggregate data volumes that must be filtered in

real-time to manageable levels for storage and analysis. The WLCG infrastructure maintains 1.5 exabytes of distributed storage to manage historical and current data.

Velocity: The velocity challenge is extreme with 40 million bunch crossings per second generating continuous data streams. The Level-1 trigger must make decisions in 3.8 microseconds. The High-Level Trigger processes 110 kilohertz of events continuously, averaging 450 milliseconds per event while maintaining a sustained 9 gigabytes per second data throughput after Level-1 filtering. Events rejected by the HLT are permanently lost and cannot be recovered, creating an irreversible velocity constraint.

Variety: The CMS detector produces heterogeneous data from multiple detector subsystems operating simultaneously. Data includes integer detector coordinates from the silicon tracker, floating-point energy measurements from electromagnetic and hadronic calorimeters, muon system information from different technologies, and higher-level reconstructed objects such as track candidates and energy clusters. Different trigger algorithms interpret this diverse data differently, and the system must accommodate physics searches as diverse as Higgs boson production at 1 hertz rates, top quark events at 20 hertz, W boson production at 4 kilohertz, and b-quark jets at 10 megahertz.

Veracity: The CMS trigger must maintain high physics accuracy while making rapid automated decisions. Track reconstruction achieves 85-95% efficiency despite detector issues including 5% disabled pixel modules. The trigger selection must balance selecting potentially important physics events against a background of increasingly common pile-up interactions that can mimic signal signatures. Machine learning algorithms incorporated into modern triggers are continuously validated against physics benchmarks and cross-checked against multiple independent reconstruction approaches to ensure reliability. The irreversible nature of trigger decisions means accuracy requirements are absolute—events lost to poor trigger decisions represent lost science that cannot be recovered in later analysis.

Value: The value is incalculable in terms of fundamental scientific discovery. The CMS trigger enabled detection of the Higgs boson in 2012, confirming a 40-year theoretical prediction. The system continues to enable searches for physics beyond the Standard Model and precision measurements of known particles. The LHC collision data requires analysis of hundreds of petabytes to extract fundamental physics, and the trigger system determines which events even have the opportunity to be analyzed. Without the HLT achieving 10 orders of magnitude data reduction while preserving scientifically

valuable events, the physics output of the multi-billion dollar facility would be impossible.

Primary Citations

CERN, "Description & performance of the current CMS trigger" (2024). Presentation at "Triggering Discoveries in High Energy Physics III" conference by M. Musich on behalf of the CMS Collaboration. Available at
https://indico.cern.ch/event/1361472/contributions/5972904/attachments/2984490/5256753/MM_HighLevelTrigger_TDIII.pdf

arXiv:2410.17038v2, "Performance of the CMS high-level trigger during LHC Run 2," November 2024

CERN, "Large Hadron Collider - Computing," <https://home.cern> (archived materials)

Worldwide LHC Computing Grid Documentation, "WLCG Infrastructure and Performance," as of 2024

CMS Collaboration, "Local reconstruction software for the CMS silicon strip tracker" and related technical documentation

CERN Courier, "The CMS experiment puts physics onto the menu" (2022)

Helmholtz Association, "As much data as Facebook," Big Data at CERN interview with Prof. Joachim Mnich (2014)